

INTRO




## INTRO

- Welcome to BIOF 510! This course will be a continuation of BIOF 509
- My info: [anibaljt@nih.gov](mailto:anibaljt@nih.gov)
- TA info: [parkyc@nih.gov](mailto:parkyc@nih.gov)

# THE COURSE

- We will cover many of the same algorithms as 509 (and some new ones) in more advanced contexts
- Course will consist of two projects, a literature review, and participation (in slack)
- Live lectures on Monday, these will be recorded and posted on canvas
- Christina will post supplementary resources midweek

# REVIEW OF DEEP LEARNING




# INPUT LAYER?



What is the input layer?



Whatever you are wanting to analyze!




## HIDDEN LAYERS

The hidden layers are the weighted combinations of features

Weights are initialized randomly and optimized






**Input  
Layer**

**Hidden  
Layer**

**Output  
Layer**



## OUTPUT LAYER




In most cases, there one node per label you can predict



Features strongly correlated to a label will have high weights for that corresponding neuron



Label corresponding to the neuron with the highest value is assigned to the input data point




# NEURAL NETWORK COMPONENTS

Now, we will cover  
what make neural  
networks so  
powerful

I. Activation  
Function

II. Loss Function

III. Optimization  
Function




Sometimes data  
cannot be separated  
using linear methods!



Activation functions  
transform the value  
of each neuron into a  
non-linear space

## ACTIVATION FUNCTIONS

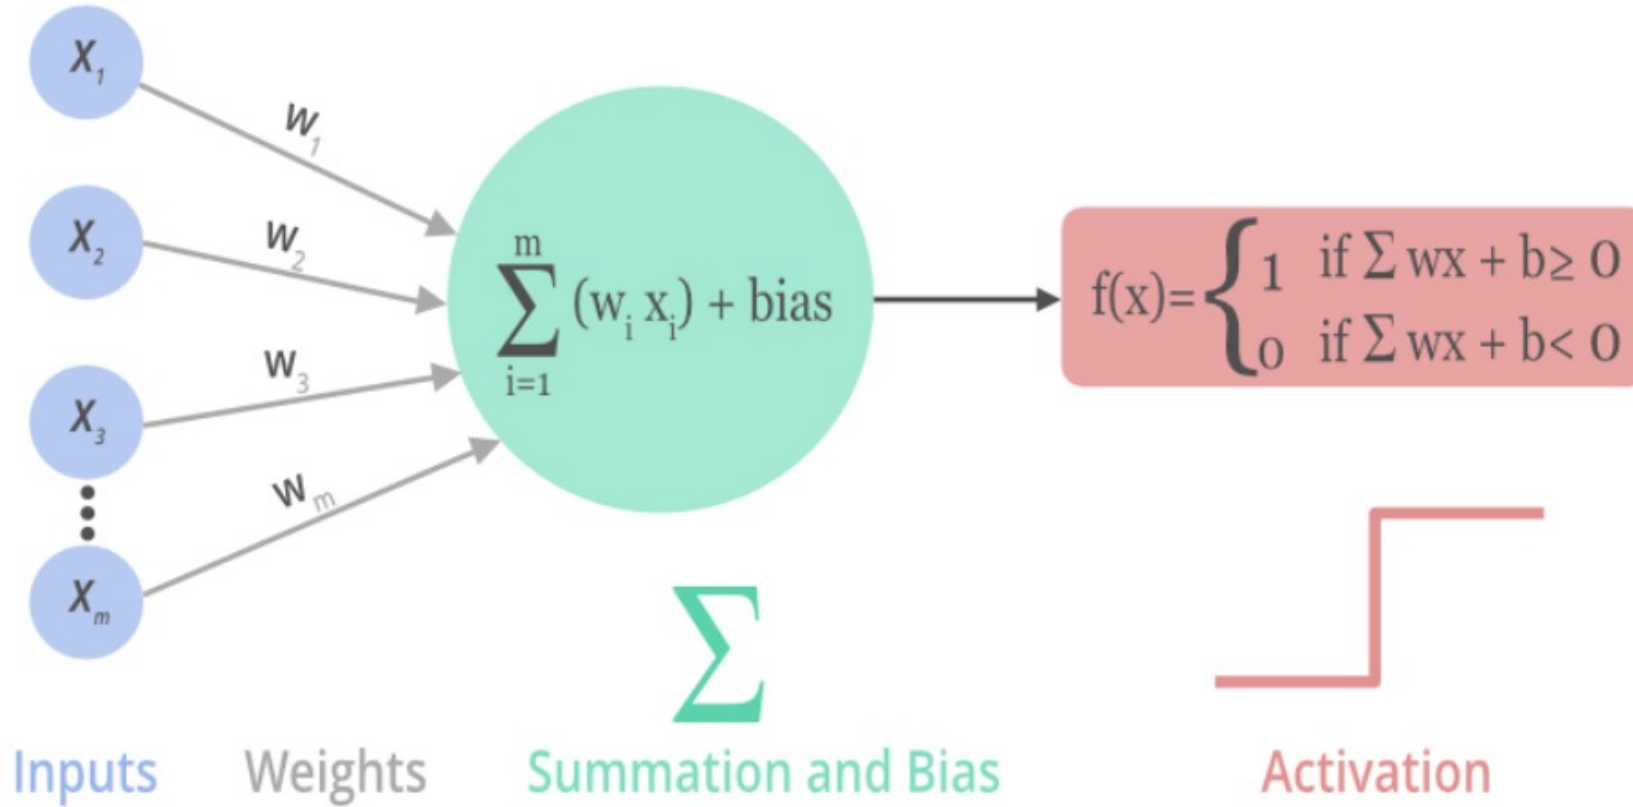



# ACTIVATION FUNCTIONS

- The rectified linear unit function (ReLU) is the most widely used activation function
- ReLU does not affect values which are greater than zero, sets all negative values to zero
- This prevent negative neurons (mostly negative weights) from being inputs into future neurons
- These “bad” features would hinder learning + classification by canceling out good features









ReLU Activation Function

## LOSS FUNCTIONS

---

We know that neural networks update the weights based on prediction accuracy

---

How does the neural network assess prediction accuracy? With a loss function

---

Labeled data has absolute ground truth (so the probability value = one)

---

Output layer of the network calculates the probability that a data point has a given label

---

Most basic way to calculate error is through subtraction

---

Data point is in class 2, NN thinks data has a 73% probability of being in class 2:  $1 - .73 = 0.27$

## LOSS FUNCTIONS

Loss functions are cleverer than just subtraction

We want wrong, confident predictions to have a large loss (weight adjustments will be bigger)


Cross-entropy loss is a very common choice for machine learning tasks

Can accommodate binary and multiclass data

## LOSS FUNCTIONS


- Cross entropy loss uses the log of the probability to severely penalize confident mistakes
- In multiclass cases, adds up the loss from across all the classes (see figure)
- In multiclass case, several big mistakes are worse than a bunch of small ones

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$



# OPTIMIZATION FUNCTIONS

- The error is derived from the loss function, and we want to move down the slope of the error function to find the minima
- Partial derivatives are used to determine the slope of the error function for each feature/weight pair
- Once we have the individual slopes, increase/decrease the weights so the neurons will contribute less to the error
- Excellent example (with actual numbers) [here](#)



# CONVOLUTIONAL NEURAL NETWORKS



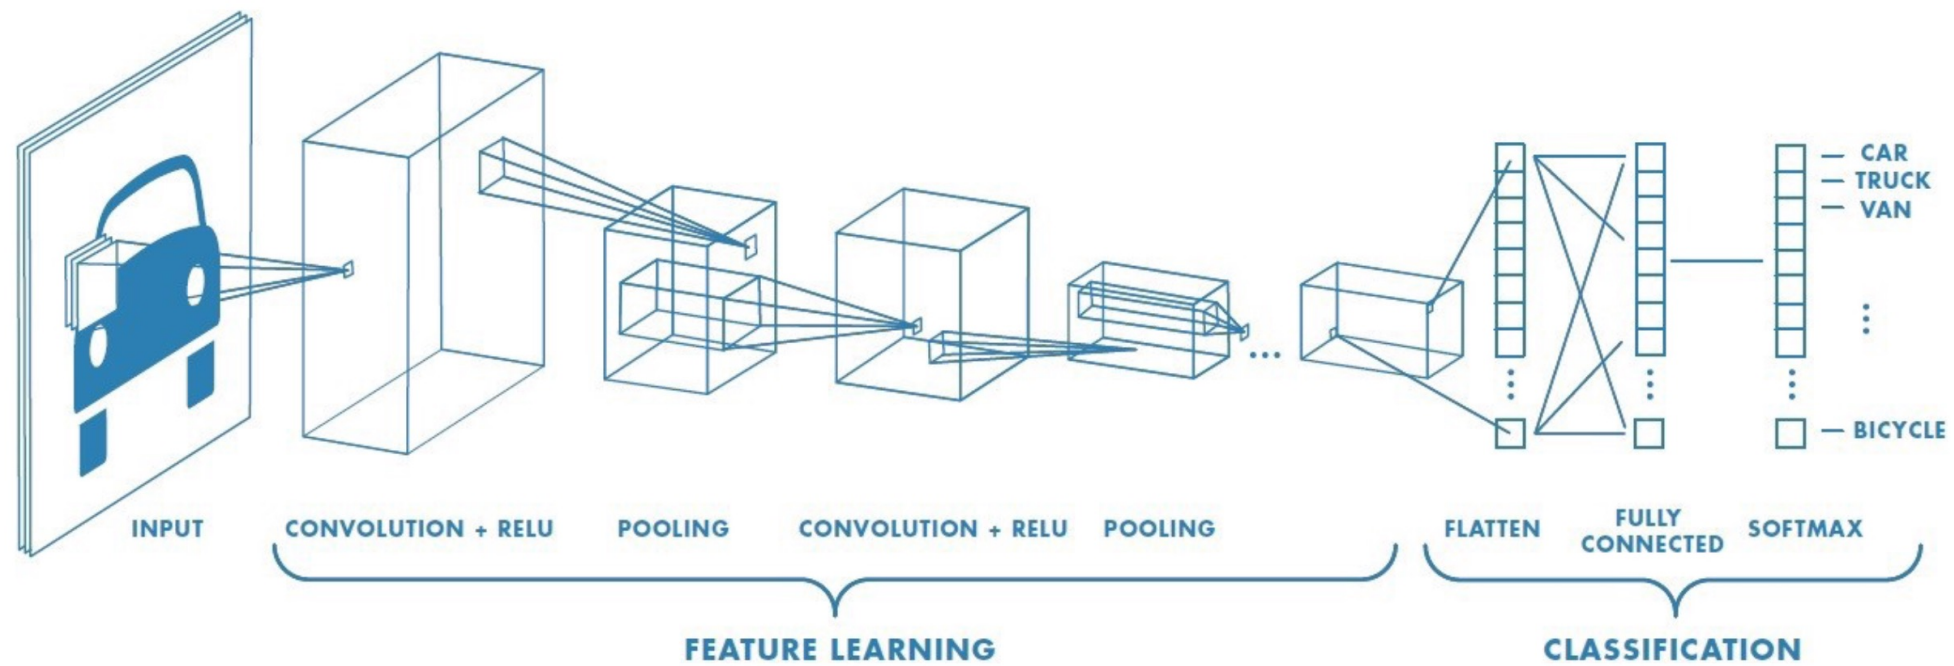
## CONVOLUTIONAL NEURAL NETWORK (CNN)

- CNNs combine the convolution operation with the traditional ANN structure
- The convolution operation captures local details about an image/matrix
- Used in image analysis to identify important features that can be used for classification

## CONVOLUTIONAL NEURAL NETWORK (CNN)

- Many different convolution kernels are passed over the regions of the data
- This captures different perspectives of each region that could represent key features
- These features are fed into an MLP to optimally weighted feature combos





# MAX POOLING



Selects the highest kernel output value in a localized region of the data




This value is used to represent whatever structure/object/entity is in that local region

12	20	30	0
8	12	2	0
34	70	37	4
112	100	25	12

$2 \times 2$  Max-Pool

20	30
112	37



AUTOENCODER

## THE AUTOENCODER


A common type of “unsupervised” deep learning is dimensionality reduction via the autoencoder

Strip the data down to a few features which correspond to combinations of features transformed by an activation function


Optimize the weights of these features such that the output layer is identical to the input layer

The resultant features capture the useful info in the data and can be used for low-dimensional, noiseless analysis (aka clustering)











LSTM

# RECURRENT NEURAL NETWORKS



ANNs connected together

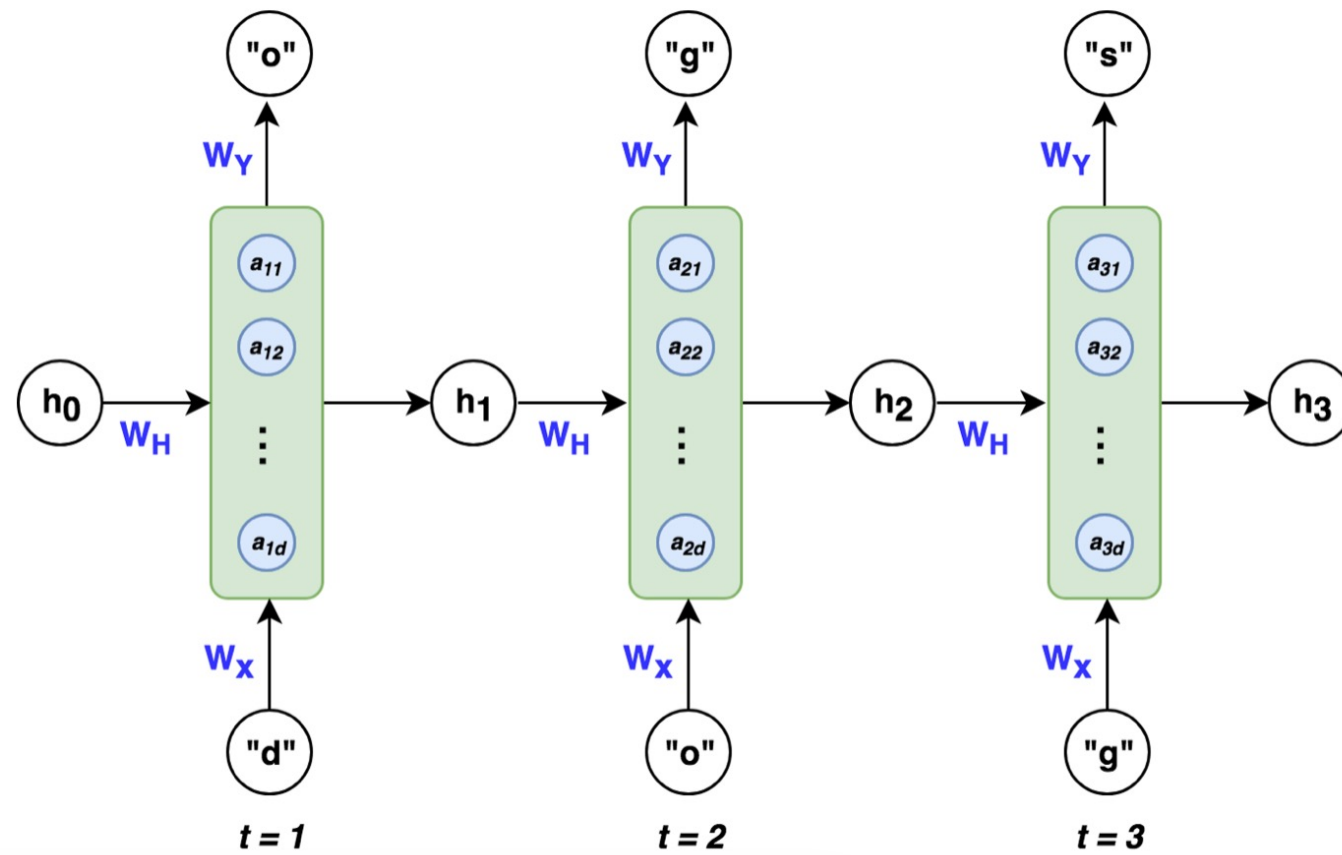


Information from prior networks are used to influence predictions of the future ones



Simple example of what RNNs can learn:  $x$  is important if  $y$  occurred in a prior state

## Simple Example



# THE EPIC LSTM



LSTMs have 3 main components



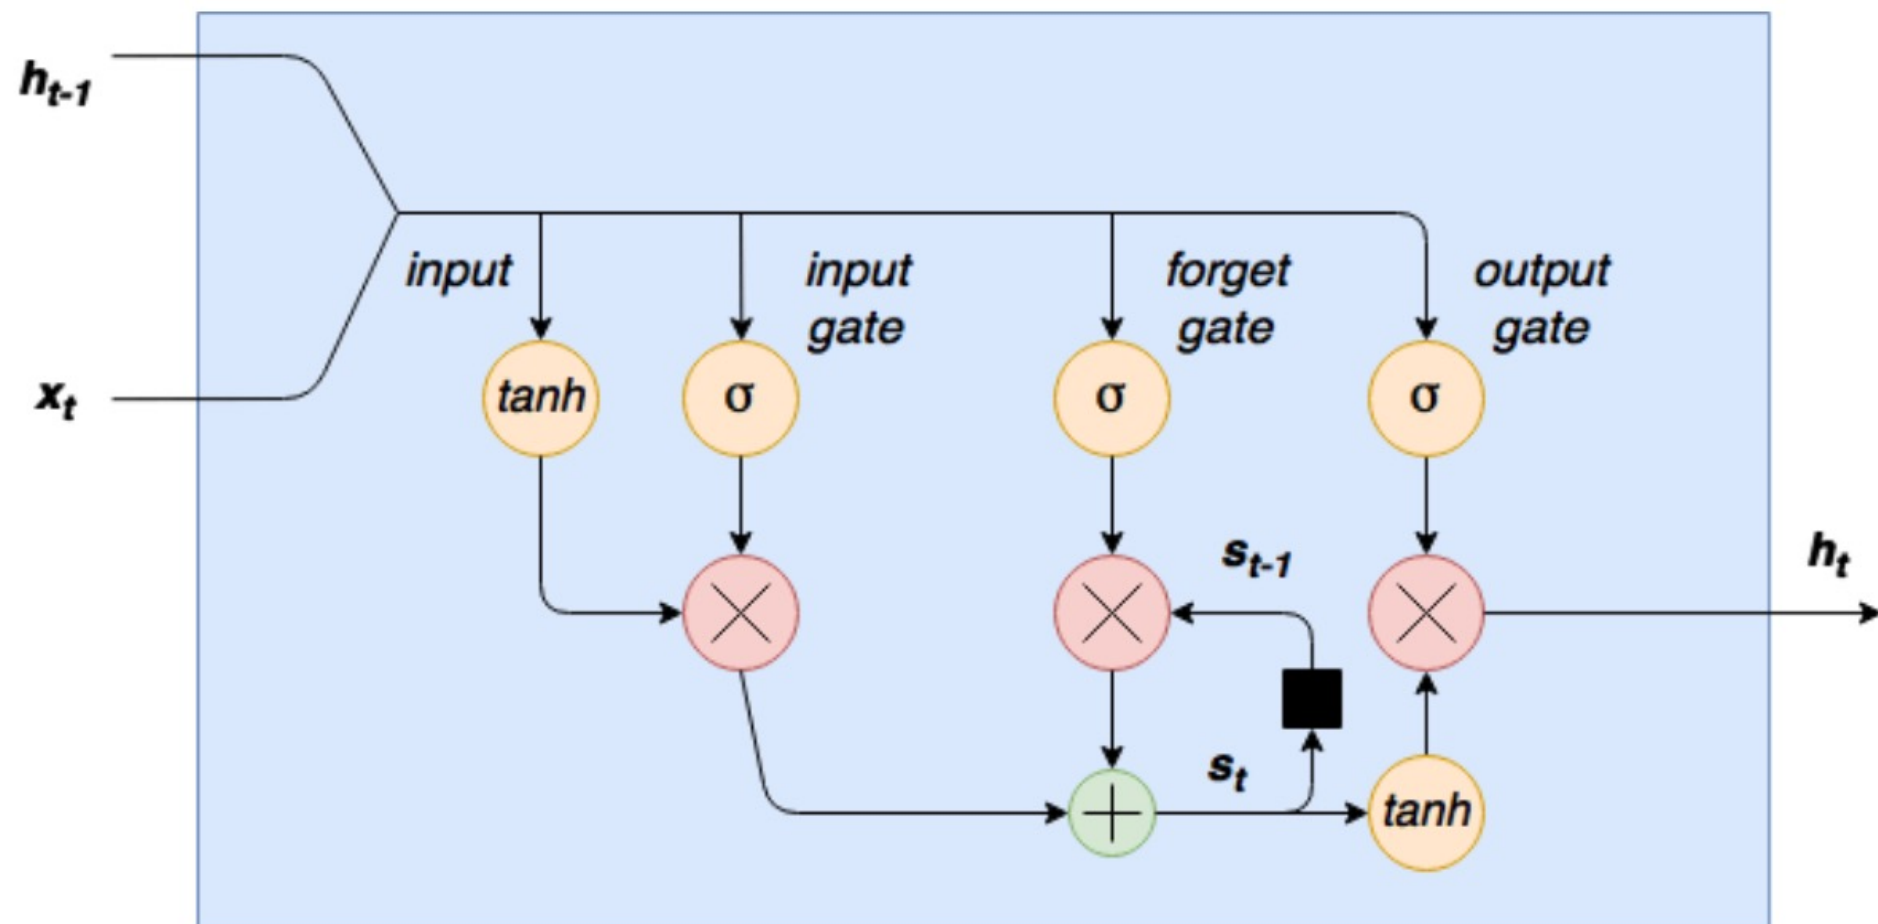
I. Long Term Memory  
(important long-term info)



II. The input from the prior state (short term info, like RNN)



III. Output – determining what moves on





# GENERATIVE ADVERSARIAL NETWORKS (GAN)

# GENERATIVE ADVERSARIAL NETWORKS



Deep Learning becomes a game

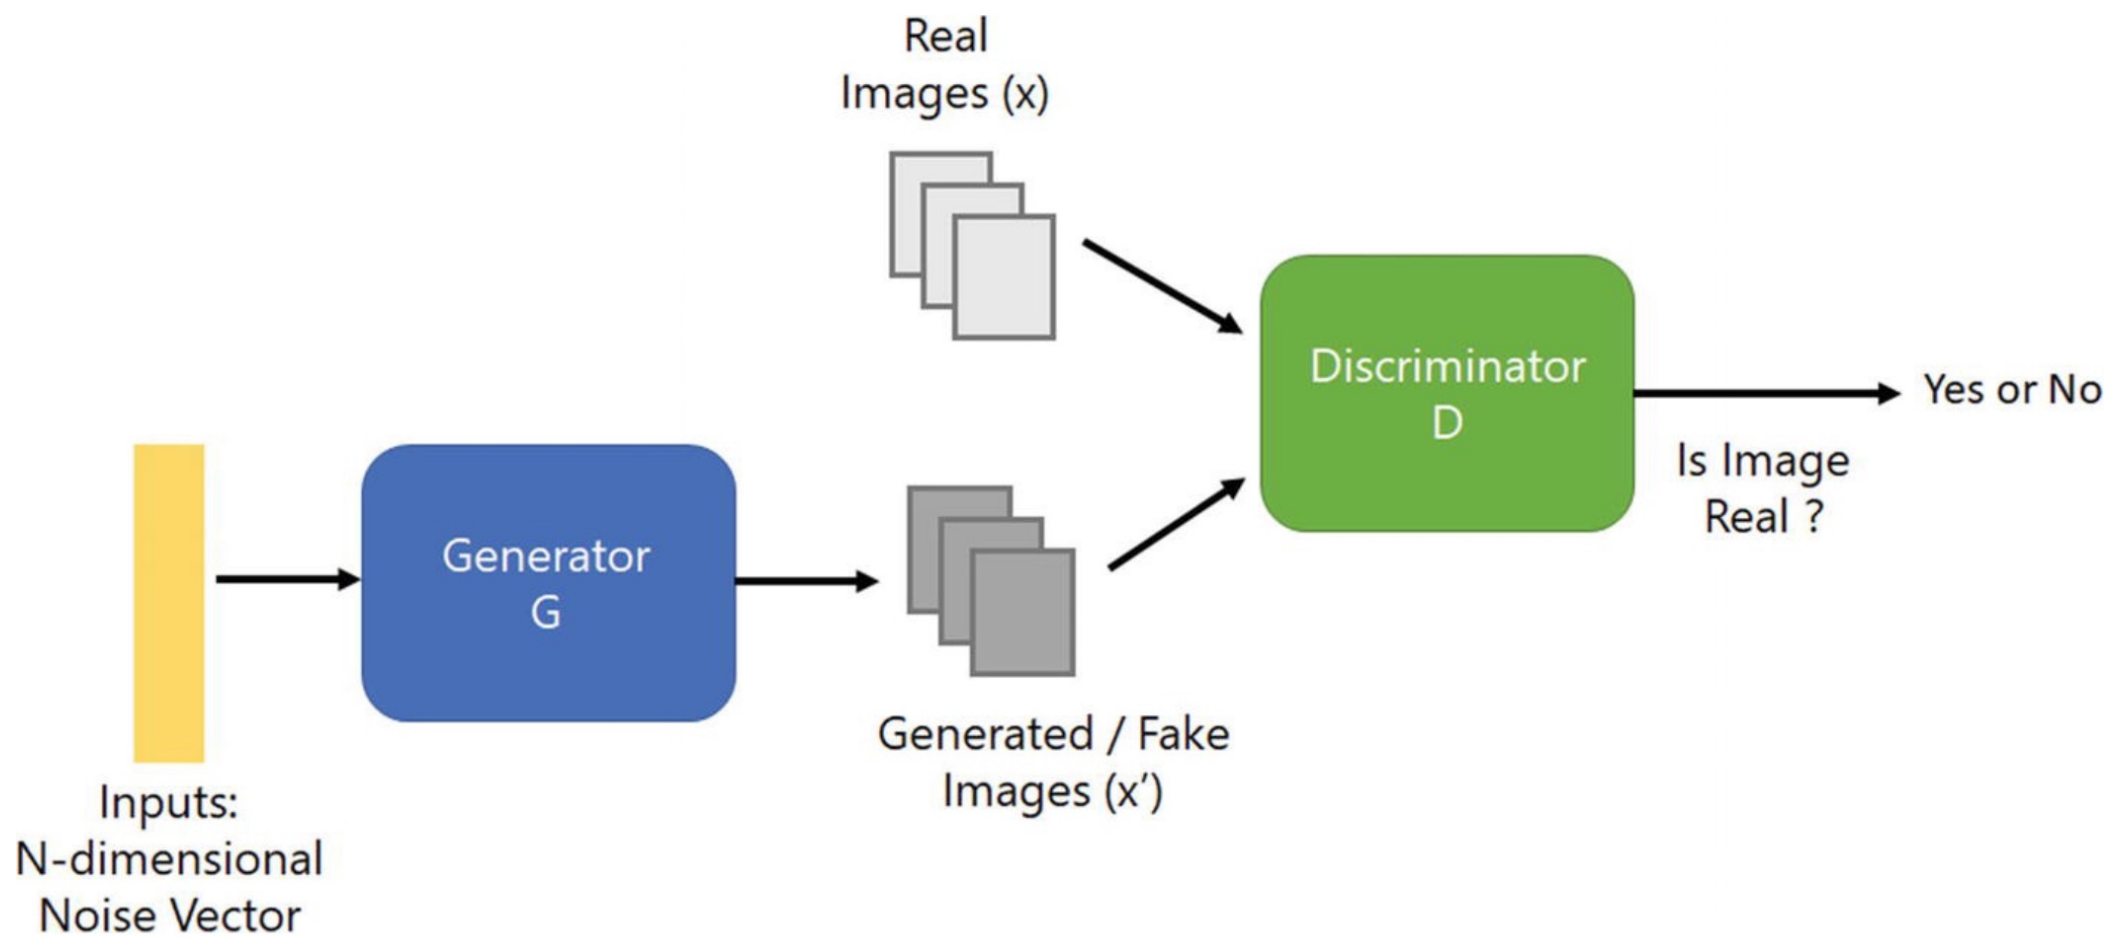


Two neural networks compete against each other




One of them generates fake data, and other tries to distinguish real from fake





FOR  
IMAGES



# TRAINING A GAN



Remember, this is a game



If the discriminator does well, the generator must be losing



If the generator does well, the discriminator must be losing



Lower loss for the discriminator = higher loss for the generator



Higher loss for the discriminator = lower loss for the generator

# TRAINING A GAN



If we train both networks at the same time, it will be like trying to hit a moving target



We train the generator for a few epochs, and do not change the weights of the discriminator



We then train the discriminator for a few epochs, do not change the weights of the generator



Repeat for a certain number of epochs until the discriminator fails half the time

## GENERAL GUIDELINES

Use the same things you would on a regular neural network

Activation functions

Optimization functions

Loss functions

Batch normalization (prevents GANs from replicating the same sample over and over)

Dropouts